

HTTP in Action

Some experiments using HTTP

Exercises

1. **Use netcat as HTTP server.** Send a request from web browser. Send HTTP response using netcat.

So you can see what a web browser really sends.

2. **Use netcat as web client.** Send a request to a real web server.

So you can see what a web server really sends.

3. **Redirect a web browser** using HTTP response codes and Location header.

Send & view HTTP using netcat

netcat (nc) or ncat are tools that let you:

- manually send and receive TCP packets
- act like a server that accepts client connections

■ "netcat" (nc) is a Linux command.

There is a version for Windows. MacOS?

■ "ncat" is newer implementation that supports IPv6.

Exercise 1: Use netcat as a server

In a terminal window. run netcat as a TCP server listening on **port** 8080

```
netcat -v -l 8080
```

- **-l** means listen for connections, **-v** means verbose
- Port 80 is the standard port for http, but you must be **root** (*admin*) to use port 80. You can use any available port 1 - 65535
- To receive a request from a **different** host, make sure there is no firewall blocking tcp port 8080 (or 80).

What is my browser sending?

Open a web browser and send a request to netcat:

`http://localhost:8080/make-my-day`

or, if netcat is using the standard http port (80):

`http://localhost/make-my-day`

you can request anything you like instead of "make-my-day".

What did the server receive?

The "netcat" console should print something like this:

```
GET /make-my-day HTTP/1.1
```

```
Host: localhost
```

```
Connection: keep-alive
```

```
User-Agent: Mozilla/5.0 (X11; Linux x86_64)
```

```
    AppleWebKit/537.36 (KHTML, like Gecko)
```

```
    Chrome/69.0.3497.81 Safari/537.36
```

```
Accept:
```

```
    text/html,application/xhtml+xml,application/xml;
```

```
    q=0.9,image/webp,image/apng,*/*;q=0.8
```

```
Accept-Encoding: gzip, deflate, br
```

```
Accept-Language: en-US,en;q=0.9,th;q=0.8
```

The Browser is Waiting for a Reply

You can see the browser is waiting for a reply.

We will use netcat to send a reply.

We are a human web server!

Send a Reply using HTTP protocol

In the netcat cmd window, type a reply using HTTP.

First line **must** be "HTTP/1.1 200 OK"

The **body** of the reply is HTML.

```
HTTP/1.1 200 OK
```

```
Content-type: text/html
```

```
        <--blank line
```

```
<html><body>
```

```
<h1>Hello, Nerd</h1>
```

```
(you can send more stuff as html)
```

(more stuff and will be appended to the browser page as you send it). Press CTRL-C to stop the server.

What does a *real* web server send?

Now we know what a request from a web browser looks like.

What does a *real* reply from a *real* web server look like?

Exercise 2: Use netcat to send http request

Use netcat to send an HTTP request to `http://www.ku.ac.th/`

You have to enter the HTTP request yourself !

```
cmd> netcat -v www.ku.ac.th 80
GET / HTTP/1.1
Host: www.ku.ac.th
(enter a blank line)
```

- -v means verbose. Netcat will print when you are connected.
- Another way: `curl -v http://se.cpe.ku.ac.th/`
curl can also use [https](#).

What does server's reply mean?

What does the reply mean?

```
HTTP/1.1 301 Moved Permanently
Server: Apache
Location: https://www.ku.ac.th/
Content-Length: 229
Content-Type: text/html
... (more header and body)
```

- Status codes 301, 302, and 303 are redirects
- A web browser will automatically go to new URL.

Plain "http" sites are hard to find

Entire web is moving to https only.

But there are a few. Try: `www.rd.go.th`

```
cmd> netcat -v www.rd.go.th 80
```

```
GET / HTTP/1.1
```

```
Host: www.rd.go.th
```

What does the response say?

Whose web site is this?

Can We Redirect a Browser?

Exercise: Use netcat to redirect web requests to Facebook or wherever you want.

1. Start netcat in listening (server) mode:

```
cmd> netcat -v -l 8080
```

2. Use a web browser, goto **http://localhost:8080**

Redirect the Browser

3. Redirect the browser to Facebook (or anyplace).
Send status code **302** Moved Temporarily (**not 301**).
If you send 301[Moved Permanently] the web browser may always go to Facebook instead of localhost.

```
cmd> netcat -v -l 8080
```

```
Listening on 0.0.0.0 (family 0, port 80)
```

```
Connection from localhost 44240 received!
```

```
HTTP/1.1 302 Don't Bother Me
```

```
Location: https://facebook.com
```

```
(blank line)
```

Did the Browser obey your redirect?

Browser should follow 302 Redirect to new Location.

You can also send a **message** in the **response body**. In case the browser doesn't follow the redirect.

```
HTTP/1.1 302 Sorry, humans not allowed
Location: https://facebook.com
Content-type: text/plain
```

Only robots allowed.

Try `https://facebook.com` instead. :-)

Optional Exercises

1. Redirect a friend's web browser.
2. Redirect from inside a web page.
3. How many requests on a page?
4. View page-load statistics using Chrome or Firefox Developer Tools.
 - see how much stuff is downloaded for a single page!

Can you Redirect your Friend's Browser?

Can you get a friend to connect to your netcat server, and redirect his browser to facebook.com?

Some issues:

1. Friend needs to know your IP address.

Type `ifconfig` or `ipconfig` to view it.

2. Your TCP port must not be blocked by firewall running on your computer. Windows: use Control Panel to create an exception .

3. KUWIN may block wifi-to-wifi connections (called Wifi isolation)

Redirect inside a Web Page?

In some situations you may want to use redirect a single web page.

You can also add a delay and display a text message:

```
<html>
<head>
<meta http-equiv="refresh"
      content="5; URL='https://facebook.com' " />
</head>
<body>
<h2>You will be redirected in 5 seconds.</h2>
</body>
</html>
```

One page, many http requests

How many HTTP requests are needed to show this page?

```
<HTML>
<link rel="stylesheet" href="stylesheet.css">
<!-- Bootstrap makes my page look cool. -->
<link rel="stylesheet" href="https://maxcdn.bootstrapcdn.com/bootstrap/3.3.7/
  css/bootstrap.min.css">
<BODY>
<h1>My vacation</h1>
<p>
For vacation we went to <a
  href="http://www.unseen.com/bangkok">Bangkok</a>.
We visited <em>Wat Phra Kaew</em>, and took this photo:
<br/>
<IMG src="images/watprakaew.jpeg" alt="Wat Phra Keao"/>
```

How Many Requests to Load a Site?

Use web developer tools to see requests, size, & time.

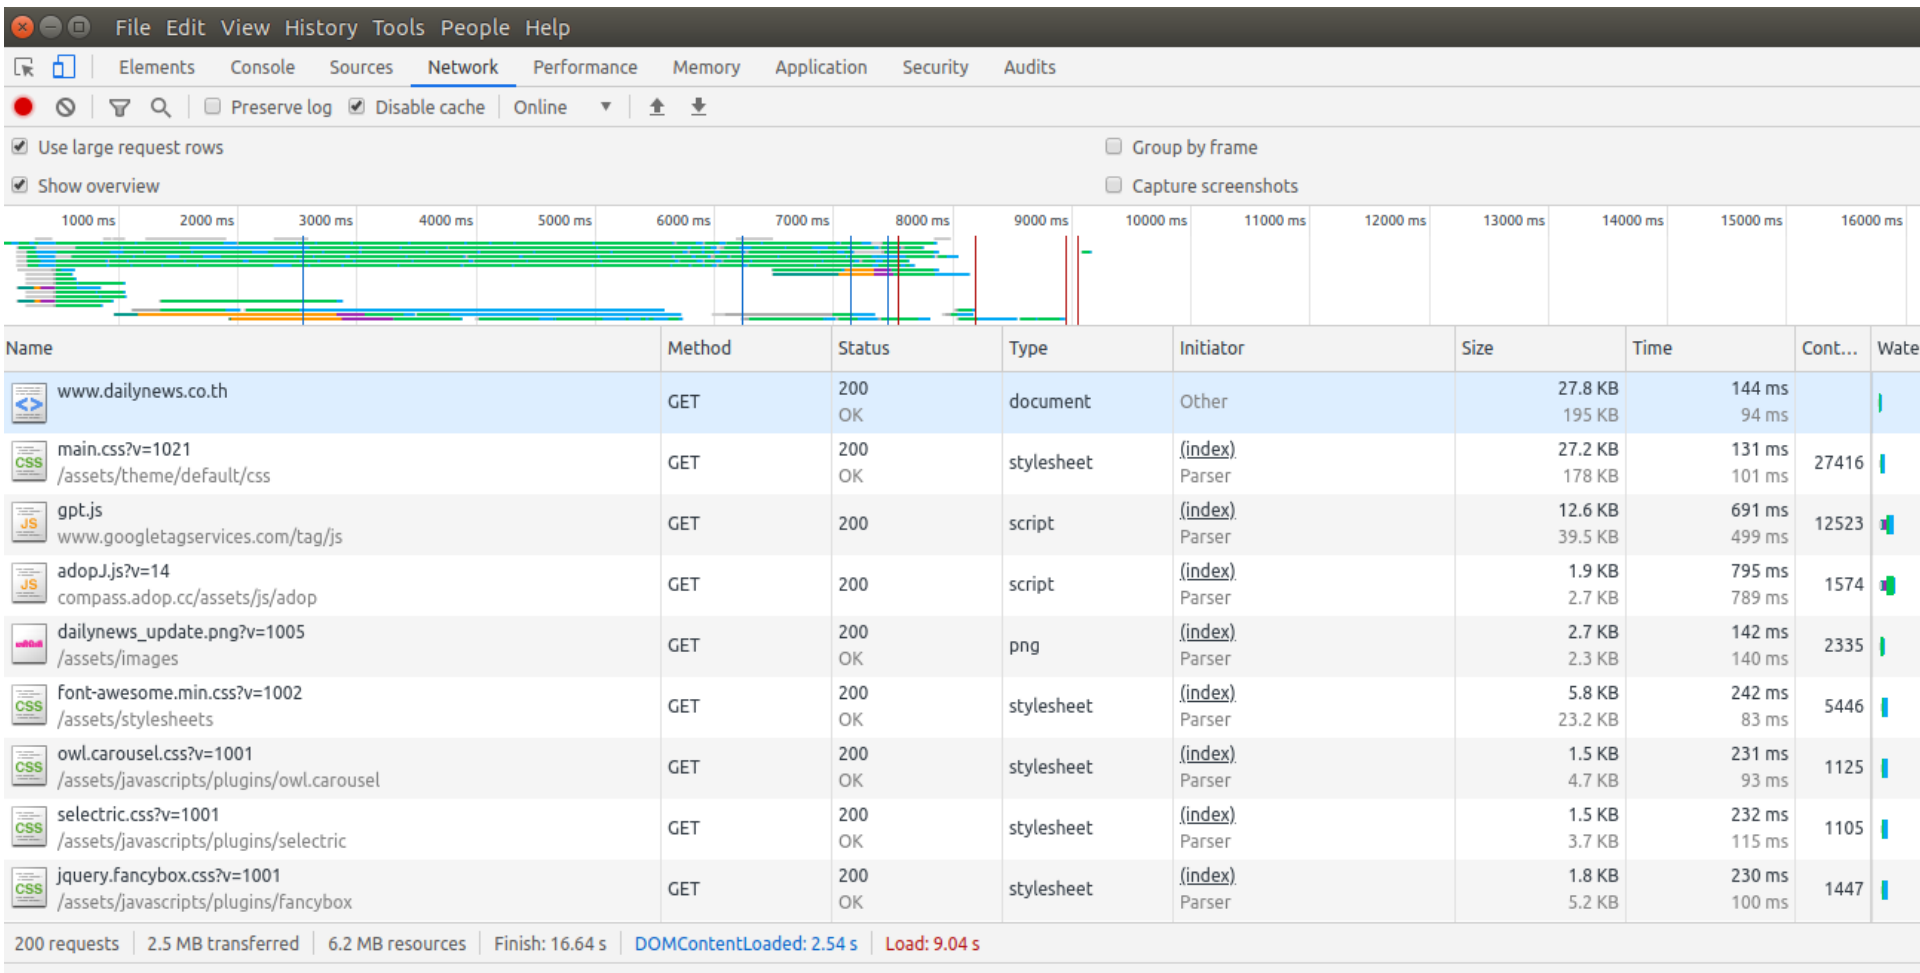
In Chrome:

1. Use *Incognito* window to avoid caching (or check a box to disable cache)
2. From "dots" menu choose **More Tools -> Developer Tools**
3. In "Developer Tools" window, choose **Network** tab.
4. In Chrome, enter a URL (such as dailynews.co.th)

How many requests? How many MB?

For just one web page!

network stats for www.dailynews.co.th



201 requests, 2.5 MB transferred, 6.2 MB resources, Load: 9.04s